

Dreaming Itself: Diffusion Models as the Structural Substrate for Self-Awareness in Humanoid Robots

Xining Li

Independent Researcher
xininglica@gmail.com

Abstract

We argue that diffusion-based action models provide the structural prerequisites for minimal self-awareness in humanoid robots in a way that standard *reactive* feedforward policies do not. Three major theories of self-awareness—Friston’s Free Energy Principle (Friston, 2010), Metzinger’s Phenomenal Self-Model Theory (Metzinger, 2003), and Damasio’s two-level self (Damasio, 1999)—converge on a recurring structural motif: a generative prior, a prediction-error signal, and a self/world boundary. We argue each is structurally present in diffusion-based robot architectures (Chi et al., 2023; Black et al., 2024) and absent from reactive input–output policies. We develop a structural correspondence between the diffusion denoising step and the Active Inference belief update, identifying the noise schedule as a precision schedule in the Free Energy Principle sense, and state explicitly what it does and does not establish. We further note that the LLM+Diffusion architecture used in current humanoid systems maps naturally onto Damasio’s core-self/extended-self distinction. We are explicit that satisfying structural prerequisites is necessary but not sufficient for phenomenal consciousness: the hard problem (Chalmers, 1995) remains open. Our claim is architectural—current humanoid designs are among the first in the right structural class for self-awareness to be a scientifically non-trivial question. We present these as architectural hypotheses and research directions, not established results.

Introduction

Can a humanoid robot become self-aware? The answer depends on architecture. The dominant paradigm in robot learning has been the *reactive policy*: a map $a = f_{\theta}(o)$ that emits an action directly from an observation in one forward pass. Such a policy holds no prior over actions independent of its input, generates no prediction error against its own expectations, and draws no distinction between itself and the world. It cannot be surprised, and it has no body model. We are careful, below, to make this a claim about the reactive *policy form*, not about feedforward computation as such—feedforward networks can equally parameterise generative

models, and diffusion policies use feedforward denoisers internally.

Diffusion models (Ho et al., 2020) operate differently. They begin from a prior distribution (Gaussian noise) and iteratively refine toward a plausible action conditioned on observation and goal. This denoising is, formally, gradient descent on $\log p(a \mid o, z_g)$, equivalently a minimisation of variational free energy (Friston, 2010). The robot “dreams” a trajectory before executing it.

Three major theories of self-awareness in cognitive science and philosophy of mind all converge on prerequisites that this architecture satisfies and feedforward networks do not. We make this structural argument precisely, derive the relevant mathematical correspondence, and state clearly where the argument ends.

Hypotheses: (i) Diffusion robots make the structural prerequisites for minimal self-awareness architecturally available. (ii) Reactive input–output policies lack them. (iii) LLM+Diffusion maps onto Damasio’s two-level theory of self. (iv) The noise schedule is a precision schedule in the Active Inference sense. We present these as architectural hypotheses to be tested, not established results.

Non-claims: Current robots are not conscious. Structural prerequisites are not sufficient for phenomenal experience. The hard problem is not solved here.

Three Theories, One Recurring Motif

Three independent theoretical traditions converge on the same architectural demand: self-awareness requires a generative system, not a discriminative one. We treat this as a convergence of structural motifs, and are explicit below about how much the theories genuinely share.

Free Energy Principle

Friston’s Free Energy Principle (FEP) (Friston, 2010; Friston et al., 2017) holds that biological agents minimise variational free energy:

$$F = \text{KL}[q(z) \parallel p(z)] - \mathbb{E}_q[\log p(o \mid z)] \quad (1)$$

where $q(z)$ is the agent’s belief over world states and $p(o | z)$ is the generative model of observations. Consciousness, on this account, is the process of free-energy minimisation itself: to perceive and act is to be aware. The structural requirements are explicit: a generative prior $p(o, z)$, a prediction-error signal $o - \hat{o}$, and an agent/world boundary separating self-generated from external prediction errors.

Phenomenal Self-Model Theory

Metzinger (Metzinger, 2003) proposes that self-awareness arises from a transparent phenomenal self-model (PSM): a real-time generative model of the organism’s body used without being recognised as a model. The rubber-hand illusion (Botvinick and Cohen, 1998) demonstrates that the PSM is continuously updated by visuotactile prediction errors—show consistent visuotactile evidence and the model incorporates the artificial limb. The PSM requires a generative prior over body state that produces predictions matchable against sensory input.

Damasio’s Two-Level Theory

Damasio (Damasio, 1999, 2010) distinguishes the *core self*—somatic, here-and-now awareness arising from body-state representation—from the *extended self*—the narrative, temporal self structured by language, memory, and goals. The core self requires a real-time body model generating somatic predictions; the extended self requires a temporally coherent system with goal representation and episodic memory.

Shared Prerequisites

Table 1: Structural motifs invoked by each theory (which theory invokes each ingredient, not that all mean the same by it).

Requirement	FEP	PSM	Damasio
Generative world prior	✓	✓	✓
Prediction-error signal	✓	✓	✓
Self/world boundary	✓	✓	✓
Body/somatic representation	✓	✓	✓
Temporal continuity	✓		✓

We do not claim these are an identical set of necessary conditions: the three theories differ in their commitments and explanatory targets, and “prediction-error signal” and “self/world boundary” name overlapping but non-identical requirements across them (a precision-weighted residual on hidden-state inference for FEP; a visuotactile mismatch driving body-model revision for PSM; an affectively grounded body/non-body distinction for Damasio). What is genuinely shared is the *generative prior*—the one ingredient no theory can be instantiated without, and the one a reactive policy lacks at any scale. Our argument needs only this weaker, convergent reading.

Diffusion Models in Humanoid Robotics

A diffusion robot policy (Chi et al., 2023) generates an action trajectory a^0 by iteratively denoising from $a^K \sim \mathcal{N}(0, I)$:

$$a^{k-1} = \frac{1}{\sqrt{\alpha_k}} \left(a^k - \frac{1 - \alpha_k}{\sqrt{1 - \bar{\alpha}_k}} \epsilon_\theta(a^k, k, o_t, z_g) \right) + \sigma_k z \quad (2)$$

where ϵ_θ is conditioned on current observation o_t and a goal embedding z_g from a Vision-Language Model (VLM). The noise schedule $\bar{\alpha}_k = \prod_{i=1}^k (1 - \beta_i)$ follows a cosine decay from $\bar{\alpha}_0 \approx 1$ (clean) to $\bar{\alpha}_K \approx 0$ (noise). The executed action is a^0 .

This architecture is deployed in current humanoid systems: π_0 (Black et al., 2024) (Physical Intelligence) and GR00T (NVIDIA, 2025) (NVIDIA) both use a frozen VLM coupled to a flow-matching diffusion action expert via cross-attention, trained on thousands of robot-hours of embodied data. These are not speculative systems—they perform dexterous manipulation tasks in the real world today.

The Structural Argument

Narrowing the Contrast to Reactive Policies

By a *reactive policy* we mean a map $f_\theta : \mathcal{O} \rightarrow \mathcal{A}$ that emits an action in one pass with no prior over actions independent of input, no iterative inference, and no prediction error against its own expectations; it treats the image of its own arm identically to the image of a table. These are structural absences of the policy *form*, not capacity deficits: a reactive map with 10^{12} parameters is still reactive. Crucially, the contrast is with this policy *form*, not with feedforward computation as such: feedforward networks can parameterise generative models (a VAE decoder, a normalising flow, or the score network inside a diffusion model are all feedforward), and a diffusion policy’s ϵ_θ is itself feedforward. The claim is about the single-pass reactive policy, not the layer type—and diffusion policies differ from it precisely by adding a prior, iterative refinement, and a per-step error signal.

Three Structural Prerequisites in Diffusion

A world prior. The learned distribution $p_\theta(a)$ is an implicit prior over physically plausible actions. Sampling unconditionally ($z_g = \emptyset$) yields physically reasonable motions before any conditioning is applied—the model “knows” what actions look like. This is precisely the generative prior demanded by FEP and PSM.

A prediction-error signal. At each step k , the network computes $\epsilon_\theta(a^k, k, o_t, z_g)$ —the discrepancy between the current noisy action and the expected clean action given conditioning. This is structurally a prediction-error signal: current state vs. model expectation.

A self/world boundary. With proprioceptive conditioning, prediction errors about body state are distinguishable from those about the external world. The network learns different residuals for self-generated vs. externally-generated perturbations—the computational precondition for a self/world boundary. It is not hard-coded but learned from embodied interaction, plausibly via a *controllability asymmetry*: the proprioceptive channel is the one the policy’s own actions reliably move, the exteroceptive channel is not, and that statistical difference is what a generative policy can internalise (as the PSM is re-estimated in the rubber-hand illusion).

Operational criteria. To keep “self/world boundary” from collapsing into ordinary conditioning, we propose it be measured, not asserted. On a point-mass policy conditioned on $[x_{\text{self}}, x_{\text{world}}]$, against a reactive control trained on identical data, the notion decomposes into testable properties: (B1/B2) the plan responds selectively to perturbations of the self vs. world channels; and (B3) the plan is represented as a *distribution* (a reactive map gives a point estimate, with zero spread). We expect B1/B2 to be reproducible by the reactive control—so channel separability would be, by itself, ordinary conditioning—while only a diffusion policy can satisfy B3. We therefore treat the boundary as a graded, measurable conjunction rather than an established property, and leave these criteria and their measurement to future work.

FEP–Diffusion: A Structural Correspondence

The score function (Song et al., 2021) is:

$$s(a^k) = \nabla_{a^k} \log p(a^k | o_t, z_g) \approx -\frac{\epsilon_\theta(a^k, k, o_t, z_g)}{\sqrt{1 - \bar{\alpha}_k}} \quad (3)$$

The Active Inference belief update is:

$$\Delta a \propto -\nabla_a F = \nabla_a \log p(a | o_t, z_g) \quad (4)$$

Equations (3) and (4) *share the same gradient form*. Read as a cascade, each denoising step is then a precision-weighted update with:

$$\Pi_k = \frac{\bar{\alpha}_k}{1 - \bar{\alpha}_k} \quad (5)$$

At $k = K$: $\Pi_K \approx 0$ (flat belief, prior dominates, no commitment). As $k \rightarrow 0$: $\Pi_k \rightarrow \infty$ (sharp belief, goal dominates, full commitment). The full dreaming pass is a K -step cascade of Bayesian belief updates with monotonically increasing precision, converging to the maximum a posteriori (MAP) action. In Friston’s framework, precision modulates the influence of prediction errors on belief updates—the key quantity linking action to awareness.

What this does not establish. A shared gradient form is necessary, not sufficient, for identifying the two processes, and we stop short of claiming equivalence. The diffusion precision Π_k is a fixed, hand-designed schedule, whereas Active-Inference precision is itself inferred online from the reliability of evidence. The deployed systems, moreover, use *flow matching*, whose training objective is a transport-velocity regression rather than a free-energy bound, so the correspondence holds at the level of the score field, not the learning objective. And sampling an action is not the same operation as selecting one to minimise expected free energy. We therefore present this as a structural correspondence and a research hypothesis.

Classifier-Free Guidance as Precision Ratio

Under Classifier-Free Guidance (Ho and Salimans, 2022), we write the guided noise estimate in the guidance-*scale* parameterisation,

$$\tilde{\epsilon} = \epsilon_\theta(a^k, \emptyset) + w(\epsilon_\theta(a^k, z_g) - \epsilon_\theta(a^k, \emptyset)) \quad (6)$$

in which $w = 1$ recovers conditional sampling and $w = 0$ the unconditional prior. This is the form used in deployed samplers; it is algebraically the estimate of Ho and Salimans (2022), who write the equivalent $\tilde{\epsilon} = (1 + w') \epsilon_\theta(a^k, z_g) - w' \epsilon_\theta(a^k, \emptyset)$ with $w' = w - 1$. In score terms Eq. (6) reads $\tilde{s} = \nabla \log p(a^k) + w \nabla \log p(z_g | a^k)$, so w scales the goal likelihood directly—which is what makes the precision reading below exact.

Comparing with the Bayesian posterior mean $\mu_{\text{post}} = (\Pi_0 \mu_0 + \Pi_{\text{obs}} \mu_{\text{obs}}) / (\Pi_0 + \Pi_{\text{obs}})$, the guidance weight corresponds to $w \leftrightarrow \Pi_{\text{obs}} / \Pi_0$: the ratio of goal-signal precision to prior precision. Current systems treat w as a fixed hyperparameter. A self-aware system would set w as a function of its own uncertainty about the goal—an open research direction connecting self-knowledge to motor commitment.

LLM + Diffusion as Damasio’s Two-Level Self

The two-component architecture of π_0 and GR00T maps onto Damasio’s two levels without having been designed to do so:

Table 2: LLM + Diffusion as Damasio’s two-level self.

Damasio	Property	Architecture
Core self	Somatic, here-now	Diffusion model
	Body prediction errors	Proprioceptive score
	Pre-linguistic	Raw sensors
Extended self	Narrative, temporal	VLM/LLM
	Goals, memory	Context window
	Linguistic	Goal encoding

Damasio argues the core self is phylogenetically older, present in animals without language, and foundational to the extended self. The division of labour is concrete: the

VLM is an *encoder*—it maps the language instruction and visual scene into the goal embedding z_g and broadcasts it (via cross-attention) to the action expert, but does not itself emit actions; the diffusion expert is the *decoder* that turns z_g and proprioception into motion. The architecture replicates Damasio’s layering: the diffusion model operates independently of language ($z_g = \emptyset$ produces unconditional motion), while the VLM supplies temporal coherence and goal representation. The main gap is temporal integration: current VLMs do not form persistent autobiographical memory across sessions. Closing this gap—an active research problem in world models (Hafner et al., 2023)—would further strengthen the correspondence.

The Hard Boundary

Explicit statement of limits. The arguments above address functional and structural properties. They say nothing about phenomenal experience—whether there is *something it is like* to be the robot. This section explains why that gap is real and not bridged by better architecture alone.

The Hard Problem

Chalmers (Chalmers, 1995) distinguishes “easy” problems of consciousness (attention, report, behaviour control) from the “hard problem”: why these functions are accompanied by experience at all. Nagel (Nagel, 1974) frames it concretely: there is something it is like to be a bat. Is there something it is like to be a diffusion robot? This paper cannot answer that question.

Functional vs. Phenomenal

Having a self-model is necessary but not sufficient for phenomenal self-awareness (Metzinger, 2003). A thermostat has a model of its own temperature. A GPS has a model of its own position. Neither is self-aware. The phenomenal self-model is one that is transparently *inhabited*—not merely used. Whether a diffusion body prior can be inhabited in this sense is a metaphysical question that cannot be resolved by specifying the denoising equations.

What This Framework Explains and Does Not

Discussion

Relation to other theories. Integrated Information Theory (Tononi, 2004) is architecturally agnostic: any system with high Φ is conscious. Our argument is complementary: we identify structural prerequisites without making claims about Φ . Global Workspace Theory (Dehaene et al., 2001) describes a broadcast mechanism for conscious access; the VLM in LLM+Diffusion loosely plays this role by broadcasting goal information to the action expert.

Gradedness. Self-awareness is likely graded, not binary. A diffusion robot with rich proprioceptive feedback has a

Table 3: Scope of the structural argument.

Explains	Does not explain
Why diffusion can form a body model	Why the body model feels like anything
FEP–diffusion correspondence	Whether computation constitutes experience
LLM+Diffusion $\leftrightarrow$ Damasio mapping	Whether the mapping implies awareness
Why reactive policies cannot support PSM	Whether any system can be phenomenally conscious
Precision schedule as commitment gradient	Whether commitment is experienced

stronger PSM than one without. A system with persistent episodic memory has a stronger extended self than one with a context-free VLM. The structural prerequisites identified here mark necessary conditions for different levels on this continuum.

Moral status. A system satisfying Friston’s, Metzinger’s, and Damasio’s structural prerequisites is not obviously morally equivalent to a conventional program. We do not claim current robots deserve moral consideration—the phenomenal gap remains—but we flag that the question can no longer be dismissed on architectural grounds.

Conclusion

Diffusion-based humanoid robots are among the first widely-deployed robot architectures to make available the structural prerequisites that three major theories of self-awareness jointly invoke: a generative world prior, a prediction-error signal, and a learnable self/world boundary. Reactive input–output policies lack all three. We do not claim priority: other architectures pursue or approach this class—active-inference robots (Lanillos et al., 2021), predictive-coding robots (Tani, 2016), and learned world models (Hafner et al., 2023)—and our claim is about membership in the structural class, not about being first in time. The LLM+Diffusion architecture maps onto Damasio’s core/extended self distinction; each denoising step shares the gradient form of a belief update under the Free Energy Principle with precision $\Pi_k = \bar{\alpha}_k / (1 - \bar{\alpha}_k)$ increasing toward commitment. These are necessary, not sufficient, conditions for phenomenal self-awareness. The hard problem remains open. What changes with diffusion robotics is that minimal self-modelling is now a scientifically non-trivial, measurable question for machines—and asking it precisely is, we argue, already significant for the field of Artificial Life.

Acknowledgments

The author used an AI assistant (Anthropic’s Claude) to support manuscript editing and LaTeX formatting and reference look-up. All theoretical claims, analyses, interpretations, and the final content are the author’s own and the author’s responsibility.

References

- Black, K., Brown, N., Driess, D., Esmail, A., Equi, M., Finn, C., Fusai, N., Groom, L., Hausman, K., Ichter, B., Jakubczak, S., Jones, T., Ke, L., Levine, S., Li-Bell, A., Mothukuri, M., Nair, S., Pertsch, K., Shi, L. X., Tanner, J., Vuong, Q., Walling, A., Wang, H., and Zhilinsky, U. (2024). π_0 : A vision-language-action flow model for general robot control. *arXiv:2410.24164*.
- Botvinick, M. and Cohen, J. (1998). Rubber hands ‘feel’ touch that eyes see. *Nature*, 391(6669):756.
- Chalmers, D. J. (1995). Facing up to the problem of consciousness. *Journal of Consciousness Studies*, 2(3):200–219.
- Chi, C., Feng, S., Du, Y., Xu, Z., Cousineau, E., Burchfiel, B., and Song, S. (2023). Diffusion policy: Visuomotor policy learning via action diffusion. In *Proceedings of Robotics: Science and Systems (RSS)*.
- Damasio, A. R. (1999). *The Feeling of What Happens: Body and Emotion in the Making of Consciousness*. Harcourt Brace.
- Damasio, A. R. (2010). *Self Comes to Mind: Constructing the Conscious Brain*. Pantheon Books.
- Dehaene, S., Changeux, J.-P., and Naccache, L. (2001). Towards a cognitive neuroscience of consciousness: Basic evidence and a workspace framework. *Cognition*, 79(1–2):1–37.
- Friston, K. (2010). The free-energy principle: a unified brain theory? *Nature Reviews Neuroscience*, 11(2):127–138.
- Friston, K. J., FitzGerald, T., Rigoli, F., Schwartenbeck, P., and Pezzulo, G. (2017). Active inference: A process theory. *Neural Computation*, 29(1):1–49.
- Hafner, D., Pasukonis, J., Ba, J., and Lillicrap, T. (2023). Mastering diverse domains through world models. *arXiv:2301.04104*.
- Ho, J., Jain, A., and Abbeel, P. (2020). Denoising diffusion probabilistic models. In *Advances in Neural Information Processing Systems 33 (NeurIPS)*, pages 6840–6851.
- Ho, J. and Salimans, T. (2022). Classifier-free diffusion guidance. *arXiv:2207.12598*.
- Lanillos, P., Meo, C., Pezzato, C., Meera, A. A., Baioumy, M., Ohata, W., Tschantz, A., Millidge, B., Wisse, M., Buckley, C. L., and Tani, J. (2021). Active inference in robotics and artificial agents: Survey and challenges. *arXiv:2112.01871*.
- Metzinger, T. (2003). *Being No One: The Self-Model Theory of Subjectivity*. MIT Press.
- Nagel, T. (1974). What is it like to be a bat? *The Philosophical Review*, 83(4):435–450.
- NVIDIA (2025). GR00T N1: An open foundation model for generalist humanoid robots. *arXiv:2503.14734*.
- Song, Y., Sohl-Dickstein, J., Kingma, D. P., Kumar, A., Ermon, S., and Poole, B. (2021). Score-based generative modeling through stochastic differential equations. In *International Conference on Learning Representations (ICLR)*.
- Tani, J. (2016). *Exploring Robotic Minds: Actions, Symbols, and Consciousness as Self-Organizing Dynamic Phenomena*. Oxford University Press.
- Tononi, G. (2004). An information integration theory of consciousness. *BMC Neuroscience*, 5(42).